

Ole Winther, Technical University of Denmark

An introduction to deep learning with PyTorch

This day will cover an introduction to deep learning including motivation, the feed forward neural network, convolutional neural networks, tricks of the trade and labs in PyTorch from <https://github.com/munkai/pytorch-tutorial>

Preparation:

Browse or read: Michael Nielsen, Neural Networks and Deep Learning, <http://neuralnetworksanddeeplearning.com/> chapter 1-4

Compute and exercises:

We will use Jupyter Notebooks to run exercises in PyTorch. To get to that you need Anaconda and PyTorch if not already properly installed

1. **Before course:** Follow the instructions in <https://github.com/DeepLearningDTU/02456-deep-learning> to install until the section called Native (no need to install TensorFlow). If you already have Anaconda installed from above, just skip this step.
2. **Before course:** Install PyTorch (<https://pytorch.org/>). If you followed instruction from above, you have 'conda' package manager and Python version 3.6 installed, you will need to select those to get the PyTorch install command.
3. **Before course:** Run exercises in https://github.com/munkai/pytorch-tutorial/tree/master/1_novice (Do not run exercises involving Cuda)
4. At the course we will run exercises from https://github.com/munkai/pytorch-tutorial/tree/master/2_intermediate